ElasticFlow: An Elastic Serverless Training Platform for Distributed Deep Learning

Diandian Gu
Peking University

Yihao Zhao
Peking University

Yinmin Zhong
Peking University

Yifan Xiong
Microsoft Research

Zhenhua Han
Microsoft Research

Peng Cheng
Microsoft Research

Fan Yang
Microsoft Research

Gang Huang
Peking University

Xin Jin
Peking University

Xuanzhe Liu
Peking University

ABSTRACT

This paper proposes ElasticFlow, an elastic serverless training platform for distributed deep learning. ElasticFlow provides a serverless interface with two distinct features: (*i*) users specify only the DNN model and hyperparameters for a job, but not the number of GPUs; (*ii*) users specify the deadline for a job, but not the amount of time to occupy GPUs. In contrast to existing server-centric platforms, ElasticFlow provides *performance guarantees* in terms of meeting deadlines and alleviates users from tedious, low-level, and manual resource management.

The characteristics of distributed training introduce two challenges. First, the training throughput scales *non-linearly* with the number of GPUs. Second, the scaling efficiency is affected by worker *placement*. To address these challenges, we propose *Minimum Satisfactory Share* to capture the resource usage of training jobs to meet deadlines, and ElasticFlow performs admission control based on it. We develop a greedy algorithm that dynamically allocates resources to admitted jobs based on diminishing returns. We apply buddy allocation to worker placement to eliminate the effect of topology. Evaluation results on a cluster of 128 GPUs show that ElasticFlow increases the number of jobs that can meet their deadlines by 1.46−7.65× compared to existing solutions.

CCS CONCEPTS

• **Computing methodologies** → **Distributed computing methodologies**; • **Computer systems organization** → **Cloud computing**.

KEYWORDS

Distributed Deep Learning, GPU Cluster, Serverless Computing, Cluster Scheduling

1 INTRODUCTION

Deep learning (DL) powers many applications and services we use every day. Training DL models is an important workload in datacenters. Most of the current efforts for DL training follow the *server-centric* style [10, 16, 18, 58], where DL developers request hardware resources in the form of machine instances (physical machines, virtual machines, or containers) to run DL training jobs. Although the server-centric approach substantially advances DL applications, it has two limitations.

First, the server-centric model is too *low-level* for DL developers. DL developers need to explicitly request hardware resources and configure machines to run their jobs. Also, DL developers face a *system-wide* problem of adapting the local batch size based on GPU memory and deciding the number of workers, as well as a *DL* problem of choosing the hyperparameters when training deep nueral network (DNN) models. This is particularly challenging for DL developers who are likely to have limited expertise in systems. Second, the server-centric model does not have the flexibility to elastically scale the resource provisioning of DL jobs to promise *performance guarantees*, i.e., guaranteeing to finish a job before a given deadline. Performance guarantees are critical for production environments that require models to be (re-)trained and onboarded in time for regular product releases [16, 34], e.g., fine-tuning BERT model with daily news to update recommendation services every day. While some DL training platforms [38, 48] are not server-centric (i.e., DL developers do not need to configure the system-wide configurations), they are not aware of DL developers' performance requirements.

Our proposal. A good platform for DL training should alleviate DL developers from low-level resource management while providing guarantees for their requirements. Therefore, we propose

ElasticFlow, an elastic serverless training platform for distributed DL. Compared to existing DL platforms [10, 16, 18, 38, 45, 48, 58], the key difference of ElasticFlow is that it exposes a high-level *deadline-driven* serverless interface to DL developers with two distinct features: (*i*) DL developers specify *only* the to-train DNN model and hyperparameters for a job, without the need to explicitly declare the required number of GPUs; (*ii*) DL developers specify the desired deadline for a job, but not the exact amount of time to occupy GPUs. This deadline-driven serverless interface decouples the DL problem from the system problem. The benefit of this design is that it alleviates DL developers from low-level tedious resource management tasks and allows them to focus on model development. Importantly, by decoupling static resource configuration from DL training code, ElasticFlow is able to leverage elastic scaling to dynamically allocate resources for every single job and provide *performance guarantees* to meet their deadlines.

Challenges. Although serverless interfaces can provide possible design space, the characteristics of DL jobs introduce two challenges to elastic resource allocation while guaranteeing given deadlines. First, the throughput of a DL job scales *non-linearly* along with the number of GPUs. The reason is that the communication overhead among workers increases with the number of workers. Second, the scaling efficiency is affected by the *placement* of workers [58]. The workers within the same server can leverage the high-bandwidth NVLink or PCIe to communicate among GPUs, while the bandwidths across servers are usually at a lower level. Thus, the placement of workers changes the scaling curve of a job. As ElasticFlow needs to consider a set of non-linear scaling curves for each job, these two challenges intertwine with each other, which further complicates the resource allocation.

Key techniques. To maximize the resource efficiency under non-linear scaling curves, we propose *Minimum Satisfactory Share*. It is known that the scaling curves of DL jobs are *concave*, which means adding resources to a distributed job may have *diminishing returns*. We use the minimum satisfactory share to capture the minimum amount of resources a job needs to meet its deadline. The key techniques of ElasticFlow include an admission control module that decides whether to admit an arriving job and a resource allocation module that dynamically allocates resources to guarantee the jobs' deadlines. For each admitted job, the resource allocation module allocates at least the minimum satisfactory share for each job to guarantee their deadlines. For the remaining resources, the module uses a greedy algorithm to prioritize resource allocation for the most efficient jobs based on the diminishing returns of their scaling curves. We prove that the algorithm is optimal under concave scaling curves. We also describe the extensions of ElasticFlow to accommodate best-effort jobs (i.e., jobs without deadlines).

We apply the buddy allocation to address the challenge of topology-dependent placements, which enables ElasticFlow to *decouple* job placement from admission control and resource allocation. In the case that jobs can be migrated and the number of workers is restricted to a power of two, buddy allocation guarantees to eliminate fragmentation , i.e., a job can always find a set of GPUs that are "close" to each other in topology as long as the number of idle GPUs is no smaller than that needed by the job.

Contributions. Our contributions are as follows.

- We propose ElasticFlow, an elastic serverless computing platform for distributed DL training. ElasticFlow provides a serverless deadline-driven interface to alleviate DL developers' resource management efforts, and exploits elastic scaling to guarantee deadlines for DL training jobs.
- We propose the metric of Minimum Satisfactory Share to capture the minimum number of GPUs a job needs to meet its deadline under non-linear scaling curves.
- We design an admission control algorithm that decides whether a job can be admitted to guarantee its deadline, and an elastic resource allocation algorithm that allocates GPUs to admitted jobs to maximize resource utilization.
- We implement a system prototype of ElasticFlow and integrate it with PyTorch. Evaluation results on a 128-GPU cluster show that ElasticFlow can substantially increase the number of jobs that can meet their deadlines by 1.46–7.65× compared to existing state-of-the-art solutions.

2 BACKGROUND AND MOTIVATION

2.1 Background

Serverless computing. Serverless computing, or Function-as-a-Service (FaaS), is an emerging paradigm to run workloads in the cloud [30]. Traditional Infrastructure-as-a-Service (IaaS) uses a server-centric model: cloud providers expose hardware resources as baremetal servers, virtual machines, or containers, where users have to figure out how many hardware resources it needs to run a certain workload. Comparatively, in serverless computing, resource management is entirely offloaded to cloud providers. Users only need to code their workloads using functions and submit functions to the serverless platform. As of now, serverless computing platforms do not have mature support for accelerators such as GPUs. A natural next step is to enable accelerator support on serverless computing platforms to power a wider family of workloads.

Distributed and elastic training. A DL training job trains a DNN model with a dataset. The job includes many iterations, each uses a batch of samples from the dataset to train the model with a forward pass and a backward pass. It is time-consuming to train DNN models, so distributed training is widely used to speed up the training process. In the data parallelism [31–33, 44] strategy of distributed training, each worker maintains a copy of the DNN model locally. For each iteration, each worker first trains the model independently, exchanges its gradients with other workers to aggregate the gradients, and then updates its local copy of the model and begins the next iteration. The batch size of each worker is called local batch size, and the sum of the batch size of all workers is called global batch size. There are also other strategies, such as model parallelism [12], hybrid parallelism [28], and pipeline parallelism [7, 22].

Many solutions have been proposed to optimize single-device training [27, 52, 53, 63], the communication between workers in distributed training [11, 15, 20, 24, 29, 55, 60, 61], and distributed training algorithms [35, 36]. Several efforts have explored elastic training [5, 19, 23, 39, 41, 42, 46, 47, 57, 59]. Based on these advancements, applying elastic scaling in distributed training platforms has become viable. Note that enabling elasticity to speed up a single

training job is not the focus of ElasticFlow. Instead, ElasticFlow focuses on scheduling multiple jobs in the cloud and exploits elasticity to provide performance guarantees in a serverless manner.

2.2 Limitations of Existing Solutions

Early efforts used cluster managers like Kubernetes or YARN to schedule DL jobs in the cloud without considering the characteristics of DL jobs, which results in low performance [9, 25, 26]. Recent efforts proposed specialized cluster schedulers for DL training jobs [10, 16, 18, 38, 45, 48, 58]. These efforts either followed the *server-centric* model or were regardless of DL developers' performance requirements, which have the following two limitations.

First, the server-centric model is too *low-level* for DL developers: it tightly couples a DL problem with a system problem. DL developers need to explicitly request hardware resources and configure machines to run their jobs. While containers simplify system configurations and make programs more portable, DL developers are still responsible for wrapping their programs in containers with low-level system configurations. More importantly, DL training jobs are constrained by GPU hardware resources. DL developers face a *system* problem of adapting the local batch size based on GPU memory and deciding the number of workers, which affects both the global batch size and the training throughput. Moreover, DL developers face a *DL* problem of choosing the hyperparameters when training DNN models (e.g., global batch size, learning rate, etc.). The low-level server-centric interface conflates the DL problem with the system problem. This is particularly challenging for DL developers that normally do not have expertise in systems. Although systems such as Amazon SageMaker [3] provide managed service for DL training, DL developers are still faced with both the *DL* problem and the *system* problem of hardware resource configurations.

Second, the existing solutions do not have the flexibility to elastically scale the resources of DL jobs to provide *performance guarantees* (i.e., guarantee to finish a job before a particular deadline). Most existing solutions focus on optimizing job completion time (JCT) [18, 38, 58]. While this is meaningful for many scenarios, there is another important class of scenarios where DL developers require *performance guarantees* when they have an explicit expectation of their jobs' deadline [16, 34]. For example, some production environments require models to be (re-)trained and onboarded in time for regular product releases. While some recent work [16] made attempts to consider deadlines, it still adopts a server-centric approach that lacks the flexibility to elastically scale up or down the resources of a job to optimize cluster-wide resource utilization and meet deadlines.

3 ELASTICFLOW OVERVIEW

3.1 Architecture

ElasticFlow exposes a high-level serverless interface to DL developers, based on which DL developers submit jobs to ElasticFlow. Then, ElasticFlow exploits resource elasticity to dynamically allocate resources to jobs based on their deadlines and the cluster status.
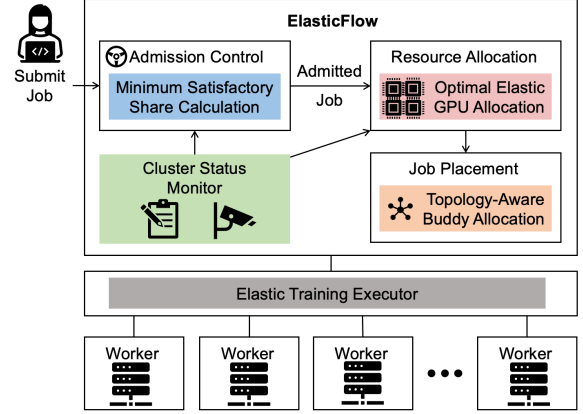

Figure 1: ElasticFlow architecture.

ElasticFlow interface. DL developers submit their training jobs as serverless functions to ElasticFlow. A function of a training job includes the following parameters.

- *DNN model*, which is the DNN model to be trained.
- *Hyperparameters*, which are the training hyperparameters such as global batch size, learning rate, etc.
- *Termination condition*, which is the condition indicating the completion of the job. DL developers only need to specify a maximum number of iterations. They can also add other conditions such as reaching a certain accuracy.
- *Deadline*, which is a point in time by which the DL developer expect the training job to finish.
- Other training components (dataset, optimizer, etc.).

This serverless interface differs from today's server-centric interface in two aspects. First, DL developers only submit a function that encodes a job to the platform, and the platform takes care of resource management. DL developers encode a training job in a function as *single-device* training and only need to specify the global batch size and other hyperparameters. This allows DL developers to focus on solving the DL problem, i.e., tuning the *global* batch size and the learning rate to achieve high model accuracy. The systems problem of deciding the *local* batch size and the number of workers based on the GPU memory is handled by ElasticFlow. Second, DL developers specify a deadline for each job. This simplifies the interaction between DL developers and the platform, as DL developers no longer need to control when to terminate allocated machine instances. It also provides flexibility for the system to dynamically adjust the resources allocated to each job based on their deadlines and termination condition.

Note that in practice, DL developers can use a variety of conditions to decide the termination of a training job, e.g., the training accuracy is above a threshold, etc. Because of the stochastic and unpredictable nature of DL jobs (e.g., the accuracy may never exceed the threshold due to bad initialization parameters), it is common to assign a maximum number of iterations to bound the running time of a job. We use the maximum number of iterations as the main termination condition in the interface. ElasticFlow also supports scheduling jobs without deadlines (§ 4.4).

ElasticFlow architecture. Figure 1 shows the architecture of ElasticFlow. DL developers submit DL training jobs to ElasticFlow. For each incoming job, ElasticFlow first uses the admission control module to decide whether to admit or drop the job. The admission control module obtains the current cluster status from the monitor module and computes the minimum satisfactory share for the job (§ 4.1) The resource allocation module then schedules the admitted jobs to efficiently utilize resources and meet their deadlines (§ 4.2). Upon each scheduling event such as job arrival or completion, the resource allocation module may update the resource allocation of some jobs through elastic scaling, i.e., adjusting the number of GPUs allocated to a job based on the deadlines of admitted jobs and the number of available GPUs. The module also computes the local batch size (i.e., dividing global batch size by the number of GPUs) for a given job. The job placement module selects GPUs from the cluster for each job based on the topology (§ 4.3). After the placement is decided, the module sends the jobs to the elastic training executor, a plugged-in component that can be replaced by any elastic DL framework. The elastic training executor ensures that each machine executes DL jobs correctly.

Performance guarantee. ElasticFlow provides the following performance guarantee: when a DL training job is admitted to the system by ElasticFlow, ElasticFlow guarantees that the deadline of the job can be satisfied.

3.2 Challenges

Non-linear scaling. Distributed training improves the job throughput with more workers. However, the throughput does not increase linearly with the number of workers due to parameter synchronizations and other overheads. The scaling curves for DL jobs are typically concave. Figure 2(a) shows the normalized scaling curves of six DNN models measured on varying numbers of GPUs. Each machine has eight NVIDIA A100 GPUs and eight NVIDIA Mellanox HDR InfiniBand HCAs, interconnected by 200 GB/s InfiniBand network. We observe concave scaling curves from the collected throughputs. For example, the training throughput of VGG16 using a global batch size of 256 with 8 GPUs is expected to be 8× that of with 1 GPU under linear scaling. While the actual throughput with 8 GPU is only 76.07% of that with linear scaling.

ElasticFlow's deadline-aware scheduling should take the non-linear scaling into account. Deadline-aware scheduling is not a new problem: traditionally, Earliest-Deadline-First (EDF) is known to be effective for meeting deadlines [17]. The canonical setup of EDF assumes a single worker per job, and jobs are ordered by their deadlines. A straightforward solution to apply EDF to schedule DL jobs is to view the entire cluster as a logical worker, the throughput of which is the sum of the throughput of all the machines. However, this solution does not work for non-linear scaling jobs. For example, consider jobs A and B, both having the same scaling curve as shown in Figure 3(a): the throughput is 1 unit with 1 worker, and 1.5 units with 2 workers. Let the deadlines of A and B be at time unit 3 and 3.5, respectively. Let the job sizes of both A and B be 3 units of iterations. EDF first runs A and then runs B, A's deadline is satisfied, but B's deadline is violated (Figure 3(b)). Alternatively, if we use one worker for each job, the deadlines of both jobs are satisfied (Figure 3(c)).

(a) Different DNN models.   (b) Different placements.
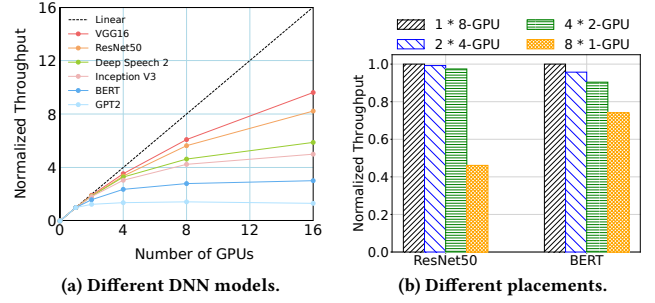
Figure 2: Characteristics of distributed training jobs. (a) Scaling curves of popular DNN models. (b) Throughputs of different placements for 8-GPU jobs.

Topology-dependent placement. The performance of a distributed training job not only depends on the number of workers, but also the placement of the workers on topology. This is because the speed of parameter synchronization between workers is decided by the communication bandwidth, which changes based on how the workers are connected. When the workers are on the same server, the communication can use PCIe (32 GB/s total bandwidth for PCIe 4.0 x16) or NVLink (600 GB/s bidirectional bandwidth for third-generation NVIDIA NVLink). When the workers are on different servers in the same rack, the bandwidth is determined by networking, which is typically 40 Gbps Ethernet, 200 Gbps InfiniBand, or 8x 200 Gbps InfiniBand for high-end GPU servers. The bandwidth is lower when the workers are in different racks.

To illustrate the problem, we measure the throughputs of training ResNet50 and BERT with different worker placements, respectively. The global batch size is 256. Each machine is configured with eight NVIDIA A100 GPUs, which are intra-connected by NVLink with 600 GB/s bidirectional bandwidth and inter-connected by InfiniBand network with 200 GB/s bandwidth. As shown in Figure 2(b), the throughputs are different depending on the placement of the workers. When the eight workers are on the same server, the throughput is 2.17× of that when the eight workers are on eight different servers for ResNet50. We only show the scaling throughputs for two models and four placements. There are more combinations, as eight workers can be divided between servers in many different ways.

The topology-dependent placement further complicates the resource allocation problem in ElasticFlow. When allocating resources to jobs, the scheduler needs to consider a set of scaling curves instead of a single scaling curve; the scaling curves also depend on the job placement.

4 ELASTICFLOW DESIGN

In this section, we describe the detailed design of ElasticFlow to address the above challenges. We propose Minimum Satisfactory Share to capture the minimum resource usage of DL jobs to meet deadlines for admission control. ElasticFlow applies buddy allocation to eliminate the effect of topology in job placement, thus reducing the number of the scaling curves to be considered for each job. ElasticFlow dynamically allocates resources to admitted jobs to maximize resource efficiency, taking the non-linear scaling into account.

(a) Scaling curve of job A and B.

(b) Earlist-Deadline-First (EDF).

(c) Optimal schedule.
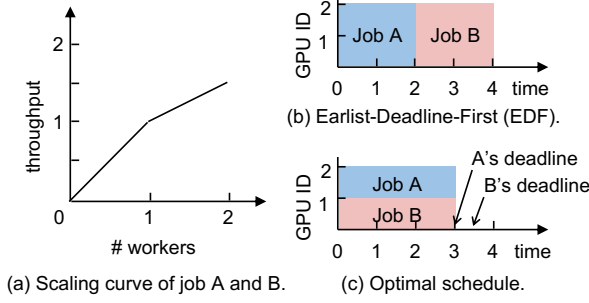
Figure 3: Motivating example to show that EDF does not work well for jobs with non-linear scaling curves.

4.1 Admission Control

Minimum Satisfactory Share. The scaling curves for DL training jobs are concave. This implies diminishing returns: the benefit of adding an extra GPU to a training job decreases with the number of GPUs. As a result, the per-GPU throughput drops as the number of GPUs increases and using a single GPU for a training job is the most efficient. We define resource usage as the number of GPUs times the running time, i.e.,

resource usage = number of GPUs × running time

For example, suppose we have a training job with a scaling curve in Figure 4(a). The throughput is 1 unit, 1.5 units, and 2 units with one, two, and four GPUs, respectively. If it requires 1 time unit to train a job with one GPU, then it would require 2/3 time unit with two GPUs and 1/2 time unit with four GPUs to train the job. The resource usage is 1, 4/3, and 2 units of GPU time with one, two, and four GPUs, respectively. Using a single GPU has the lowest GPU resource usage. However, because jobs have deadlines, we cannot train every job with a single GPU as it may violate their deadlines.

We define *Minimum Satisfactory Share* as the least number of resources needed to train a job in order to meet its deadline. The resource usage is minimized when jobs are assigned with their minimum satisfactory shares to run. For the preceding example, if the deadline of the job is 1 time unit, then we can simply allocate one GPU to the job. But if the deadline is 2/3 time unit, then we need to allocate two GPUs to the job to meet its deadline. Allocating two GPUs minimizes resource usage while meeting the deadline of the job.

Admission control. Only jobs that ElasticFlow can guarantee their deadlines are admitted. Specifically, ElasticFlow computes the minimum satisfactory share of each arriving job. If the system has more resources than the minimum satisfactory share, the job can be admitted as its deadline can be guaranteed. Otherwise, the job is dropped to avoid violating the performance guarantee.

When there is no running job in the cluster, computing the minimum satisfactory share for an incoming job is straightforward: we only need to allocate the smallest number of GPUs so that the completion time of the job is no later than its deadline. This can be done efficiently with a binary search.

When there are already jobs admitted by the system, directly using binary search to solve the above problem is not sufficient. For example, suppose we have a job C with a scaling curve shown in Figure 4(a). Let the deadline of job C be 2 time units. When there

(a) Scaling curve of job C.

(b) Only schedule job C.

(c) Schedule job A, B and C.
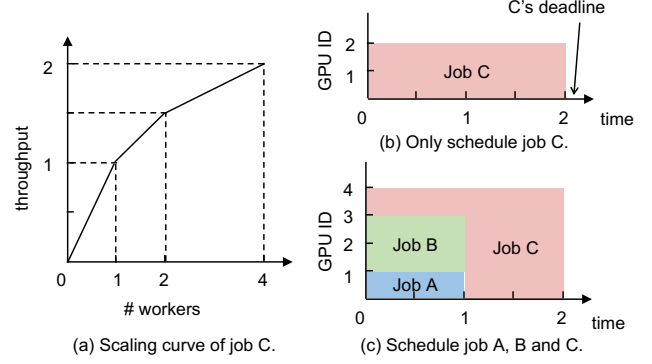
Figure 4: Example for ElasticFlow admission control. The GPU time to finish a job changes depending on other jobs. Finishing job C with the scaling curve in (a) requires 4 units of GPU time in (b) and 5 units of GPU time in (c).

is only job C, job C can use two GPUs to finish in 2 time units, as shown in Figure 4(b). Now suppose there are already two jobs A and B in the system, and suppose job A and B need to use one and two GPUs for one time unit, respectively. Because there are only four GPUs in total, job C cannot use two GPUs for the first time unit. In this scenario, the minimum satisfactory share for job C is to use one GPU for the first time unit and four GPUs for the second time unit, in order to meet its deadline.

We first describe how to solve the simple case of linear scaling curves to obtain intuitions, and then extend to non-linear curves. For the linear scaling case, the per-GPU throughput is the same regardless of the number of allocated GPUs. Formally, let $T(x_i) = k_i \cdot x_i$ be the scaling curve of job i, i.e., the throughput of job i is $T_i(x_i)$ when the number of GPUs allocated to job i is x_i and k_i is the per-GPU throughput of job i. The throughput is measured as the number of iterations per time unit. Let M_i be the maximum number of iterations to run, which represents the termination condition. Let D_i be the deadline of job i. To meet the deadline of job i, we need to allocate $M_i/(k_i D_i)$ number of GPUs to job i for D_i amount of time, or equivalently a total amount of M_i/k_i GPU time. Let G be the total number of GPUs. We have the following theorem.

THEOREM 1. *For n jobs with linear-scaling curves, we sort the jobs by their deadlines in increasing order. If the condition*

$$\forall\, i \le n - 1, \sum_{j=0}^{i} M_j/k_j \le G \cdot D_i. \tag{1}$$

holds, then there exists an allocation that can guarantee the deadlines for all jobs. Otherwise, there is no allocation that can guarantee the deadlines for all jobs.

The intuition of the theorem is that because the scaling curves are linear, it does not matter how many GPUs are used for a job in each time slot. We only need to: (*i*) sort jobs by deadlines, and (*ii*) account for the total GPU time consumed by a job and check if it exceeds the amount needed by the job before its deadline. Condition (1) ensures that the required GPU time of the jobs does not exceed the available GPU time. By ordering the jobs by their deadlines in increasing order, we can prove the theorem by induction.

Algorithm 1 Admission Control

```
1: function ADMISSIONCONTROL(job)
2:     Add job to jobs
3:     Sort jobs by deadline in increasing order
4:     for i in jobs do
5:         if PROGRESSIVEFILLING(i, 0) is False then
6:             // Drop job
7:             return False
8:     // Admit job
9:     return True
10:
11: function PROGRESSIVEFILLING(i, startTime)
12:     isSatisfied ← False
13:     for j from 1 to G do
14:         for t from D_i to startTime do
15:             x_i(t) ← min(j, G − Σ_{k=0}^{k=i−1} x_k(t))
16:             // Check if the job can be satisfied
17:             if Σ_{t=0}^{D_i} T_i(x_i(t)) ≥ M_i then
18:                 isSatisfied ← True
19:                 Break
20:         if isSatisfied is True then
21:             Break
22:     return isSatisfied
```

The problem becomes challenging for DL jobs due to their non-linear scaling curves. In this case, we cannot simply use M_i/k_i to compute the amount of GPU time needed by job i. Let $x_i(t)$ be the number of GPUs allocated to job i at time slot t. Then the admission control is to check whether we can find an allocation $A = \{x_i(t)\}$ that satisfies the following two conditions:

$$\forall i, \sum_{t=0}^{D_i} T_i(x_i(t)) \geq M_i; \tag{2}$$

$$\forall t, \sum_{i=0}^{n-1} x_i(t) \leq G. \tag{3}$$

The first condition represents that the total number of iterations performed by a job is no smaller than the maximum number of iterations of the job, i.e., the deadline is met. The second condition represents that the total number of GPUs allocated to the jobs in any time slot is no bigger than the number of available GPUs.

We apply the intuition of solving the problem under linear scaling curves and adapt progressive filling [8] to compute the minimum satisfactory share under non-linear scaling curves. The key idea is to: (i) sort the jobs by deadlines, and (ii) increase the number of GPUs for each job progressively until the deadline of the job can be met. Algorithm 1 shows the pseudocode of admission control. The algorithm first sorts the jobs by deadlines in increasing order (line 3). For each job, it checks whether the deadline of all admitted jobs can still be satisfied after adding the new job (line 4-7). It increases the number of GPUs j for job i from one (line 13). But j should not be larger than the number of remaining GPUs in the cluster (line 15). The algorithm sums up the number of iterations job i can perform at time slot t. If the sum is no smaller than the maximum number of iterations, the job can finish before its deadline (line 17-19). If job i cannot be satisfied after enumerating the number of GPUs, it means adding the new job would violate the deadline of at least one job. Therefore, the new job is dropped (line 5-7). Otherwise, it is safe to admit the new job (line 8-9).

We use the example in Figure 4 to illustrate how the algorithm works. Let $D_c = 2$ and $M_c = 3$. Suppose there are already two jobs A and B in the system and they need to use 3 GPUs for the first time slot. If j is 2, then job C can use only 1 GPU for the first time slot and 2 GPUs for the second time slot. In this case, $\sum_{t=0}^{2} T_c(x_c(t)) = T_c(1) + T_c(2) = 2.5 < M_c$, which means job C does not finish in 2 time slots. If j is 4, then job C gets 1 GPU for the first time slot and 4 GPUs for the second time slot, and we have $\sum_{t=0}^{2} T_c(x_c(t)) = T_c(1) + T_c(4) = 3 \geq M_c$, which means the deadline of job C can be met.

4.2 Resource Allocation

Allocating the minimum satisfactory share to each job can already provide performance guarantee, but after each job is allocated with its minimum satisfactory share, the cluster may still have idle GPUs left. It is beneficial to allocate the remaining GPUs to the admitted jobs, because if allocating more GPUs further speeds up the jobs, ElasticFlow would free up more GPUs even before the deadlines of the admitted jobs. This would allow the system to accommodate more jobs in the future. Our goal is to find an *efficient* resource allocation that fully utilizes the GPUs while guaranteeing the deadlines of all admitted jobs.

A strawman solution is to allocate all the remaining GPUs to the job with the earliest deadline. But this does not make the best use of GPU resources, because the marginal return of adding one extra GPU diminishes with more GPUs under non-linear scaling curves. Another strawman solution is to allocate the remaining GPUs evenly to the admitted jobs. This solution avoids adding all resources to one job, but it does not consider the current number of GPUs a job has, which affects the benefit of adding one extra GPU to a job.

From the strawman solutions, we can see that the key to achieve our goal is to account for the diminishing returns of each job given its allocated GPUs. To derive the optimal solution, we formally formulate the problem as follows. We consider the next time slot for resource allocation. Note that it is unnecessary to consider further time slots, because even if we do, new jobs may arrive and be admitted, and we would have to recompute the resource allocation plan. We use $a_i(t)$ to denote the GPU allocation that has already been allocated to job i over time. When the resource allocation starts, it is equal to the number of GPUs allocated to job i calculated by admission control (Algorithm 1). Let $x_i(t)$ be the actual number of GPUs to allocate to job i over time. Then resource allocation is to solve the following optimization problem:

$$\min \sum_{i=0}^{n-1} \sum_{t=0}^{D_i} x_i(t) \tag{4}$$

$$s.t. \ \forall i, \sum_{t=0}^{D_i} T_i(x_i(t)) \geq M_i; \tag{5}$$

$$\forall t, \sum_{i=0}^{n-1} x_i(t) \leq G; \tag{6}$$

$$\forall i, \ T_i(x_i(0) + \min(1, \ G - \sum_{i=0}^{n-1} x_i(0))) \ <= T_i(x_i(0)). \tag{7}$$

Algorithm 2 Resource Allocation

1: **function** RESOURCEALLOCATION(jobs)
2:     **for** i in $jobs$ **do**
3:         // **Allocate** $a_i(0)$ **to job** i
4:         $G \leftarrow G - a_i(0)$
5:         // **Assume job** i **gets 1 more GPU at time slot** 0
6:         // **Calculate the marginal return of job** i
7:         $x_i(0) \leftarrow a_i(0) + 1$
8:         PROGRESSIVEFILLING$(i, 1)$
9:         $i.priority \leftarrow \sum_{t=0}^{D_i} a_i(t) - \sum_{t=0}^{D_i} x_i(t)$
10:        **if** $i.finish\_time(x_i) < i.finish\_time(a_i)$ **then**
11:           Add jobs to priority queue $Queue$
12:     **while** $G > 0$ and $!Queue.empty()$ **do**
13:         // **Pick the job with largest marginal return**
14:         $i \leftarrow Queue.dequeue()$
15:         // **Assign** $x_i(t)$ **as the new allocation of job** i
16:         **for** t from 0 to D_i **do**
17:            $a_i(t) \leftarrow x_i(t)$
18:         $G \leftarrow G - 1$
19:         // **Update the marginal return of job** i
20:         $x_i(0) \leftarrow a_i(0) + 1$
21:         PROGRESSIVEFILLING$(i, 1)$
22:         $i.priority \leftarrow \sum_{t=0}^{D_i} a_i(t) - \sum_{t=0}^{D_i} x_i(t)$
23:         **if** $i.finish\_time(x_i) < i.finish\_time(a_i)$ **then**
24:            $Queue.insert(i)$

In the formulations, $\sum_{t=0}^{D_i} x_i(t)$ is the total GPU time used by job i with $x_i(t)$ GPUs at time slot t. Under sub-linear scaling, the more GPUs are allocated to a job, the more GPU time it needs to reach the termination condition. The objective is to minimize the total GPU time consumed by all jobs. There are three constraints: (i) all jobs' deadlines are all guaranteed; (ii) the total allocated GPUs do not exceed the number of available GPUs; (iii) all GPUs are allocated in the next time slot unless adding more GPUs to any job would make the job slower.

We develop a greedy algorithm to solve the problem. The intuition is to allocate the remaining GPUs to the job with the highest marginal return. Algorithm 2 shows the pseudocode. The algorithm allocates at least the minimum satisfactory share a_i to each job so that the deadlines are guaranteed (line 4) and then maintains a priority queue to order the jobs (line 5-11). The priority of a job is the marginal return to add one extra GPU to it. The algorithm greedily allocates one GPU each time (line 12-24): for each iteration, the algorithm dequeues the head from the queue (line 13-14), allocates one GPU to the job (line 18), computes the new marginal return of the job(line 18-22), and inserts the job back into the queue (line 23-24). The iterations finish until all GPUs are allocated. We have the following theorem for the optimality of the algorithm.

THEOREM 2. *Algorithm 2 computes an optimal allocation for the resource allocation problem in formula (4-7), i.e., it finds the most efficient allocation that the GPUs are fully utilized while guaranteeing the deadlines of all admitted jobs.*

The main idea of the proof of the theorem is to show that allocating each GPU to the job with the highest marginal return maximizes the utility of the GPU. It applies induction to prove the general case of multiple GPUs.

4.3 Job Placement

After the number of GPUs for a job is decided, the job needs to be placed in the cluster, i.e., selecting which GPUs to run the job. As

we have shown in §3.2, the scaling curve of a job depends on the placement. To make the problem more challenging, job placement intertwines with admission control and resource allocation, because admission control and resource allocation requires scaling curves to make decisions. If admission control uses a different scaling curve from the actual placement, the actual job throughput may be lower than the one used by admission control, leading to deadline violations.

A naive approach is to always use the most pessimistic scaling curve , i.e., when all workers of a job are on different machines. However, this potentially *underestimates* the throughput of a job, and thus *overestimates* the resource usage. This would unnecessarily reserve GPUs for admitted jobs, preventing the system from admitting more jobs.

Topology-aware job placement. We develop a topology-aware job placement algorithm to address this problem. We use a multi-layer hierarchical tree to capture the GPU topology where GPUs are connected by different types of links with different bandwidths. The leaf nodes of the tree represent GPUs. Each internal node indicates a particular connection. Figure 5 shows an example of one server with a 4-layer hierarchy. The server has two CPU sockets connected by QPI. Each CPU is connected to four GPUs via PCIe. In GPU clusters, there are multiple such servers. The servers can be connected by a Top-of-Rack (ToR) switch in the same rack, which is the 5^{th} layer in the hierarchy of the cluster.

To place a job, we find a suitable subtree that contains sufficient idle GPUs to accommodate the job. There can be multiple subtrees that satisfy this requirement. This is a bin packing problem and there are multiple heuristics. We use Best-Fit [50] which chooses the subtree of which the number of idle GPUs is closest to the number of GPUs needed. In this way, the job is allocated with the highest possible bandwidth between its workers, thus avoiding overestimating the resource usage of jobs. Therefore, job placement can be decoupled from admission control and resource allocation.

Defragmentation with buddy allocation. Our algorithm is general to jobs that require any number of GPUs, but the algorithm cannot avoid resource fragmentation. It is possible that we cannot find a subtree for a job, but the cluster contains more GPUs than needed. For example, suppose that there are two servers with the hierarchy in Figure 5. If we allocate seven GPUs of server 1 to job A and seven GPUs of server 2 to job B, we cannot find a layer-1 subtree with two GPUs for job C, although there are two idle GPUs in the cluster.

We apply buddy allocation combined with job migration used by prior DL schedulers [23, 62] to eliminate resource fragmentation. Similar to CoDDL [23], we restrict the number of workers of each job to be a power of two. Under this condition, our algorithm guarantees that there is no resource fragmentation. If the cluster has enough idle GPUs for a job, there exists a subtree that contains enough GPUs.

4.4 Discussion

Handling jobs without deadlines. ElasticFlow can be extended to support both SLO jobs (with deadlines) and best-effort jobs (without deadlines). While best-effort jobs do not have deadlines, it is desirable to finish them as soon as possible. The objective for this

Figure 5: ElasticFlow organizes GPUs in a multi-layer hierarchy, and uses buddy allocation to minimize fragmentation.

scenario is to guarantee deadlines for SLO jobs and minimize the average JCT for best-effort jobs. ElasticFlow uses a unified framework to support both kinds of jobs. It sets the deadlines of best-effort jobs to infinite and allocates GPU resources to them after allocating the minimum satisfactory shares to SLO jobs in Algorithm 2. We show the evaluation results of supporting best-effort jobs in §6.5.

Dropped jobs and hard vs. soft deadlines. ElasticFlow provides performance guarantees for hard deadlines and drops a job if its deadline cannot be satisfied. In practice, there are jobs with soft deadlines: finishing them are still meaningful even if the deadlines are missed. Supporting jobs with soft deadlines is similar to supporting best-effort jobs. ElasticFlow puts them to the queue in Algorithm 2 and allocates resources to them after allocating the minimum satisfactory shares to admitted jobs. This ensures that the deadlines of admitted jobs are always guaranteed; other jobs finish as early as possible.

Malicious users and admission control policies. We assume that users are not malicious and set their job deadlines according to their needs. It is possible that users can game the system by manipulating the job submission frequency and deadlines. For example, a user may submit many jobs with close deadlines to occupy all GPUs in the cluster, preventing the system from admitting jobs from other users. To address this problem, the cloud operator can use quotas (e.g., set a maximum number of jobs that can be submitted by each user per day) or charge users (e.g., the cost depends on the job size and the deadline). The cloud operator can apply an extra policy or charge the user before line 9 of Algorithm 1 to decide whether to actually admit the job. The policies and pricing models are orthogonal to ElasticFlow, and are interesting directions for future work.

Node failures. In real-world clusters, there might be random node failures. ElasticFlow can be extended to taking node failures into consideration by calculating the probability of node failures and reserving enough resourcse to ensure performance guarantee in case of node failures.

5 IMPLEMENTATION

We have implemented a prototype of ElasticFlow in 8,800+ LOC of Python, including 5,600+ LOC for the scheduler and 3,200+ LOC for elastic training. We integrate it with PyTorch 1.10.0 [44]. We use PyTorch DistributedDataParallel (DDP) for distributed training,

Task	Dataset	Model	Batch Size
CV	ImageNet [13]	ResNet50 [21]	64, 128, 256
		VGG16 [51]	64, 128, 256
		Inception V3 [54]	64, 128
NLP	CoLA [56]	BERT [14]	64, 128
	aclImdb V1 [37]	GPT-2 [49]	128, 256
Speech Recognition	LibriSpeech [43]	Deep Speech 2 [6]	32, 64

Table 1: DNN models used in the evaluation.

NCCL [1] for communicating between workers, and gRPC [4] to exchange control messages between the scheduler and workers.

ElasticFlow interface. ElasticFlow currently supports DL training jobs written in PyTorch. PyTorch provides an official launch API in PyTorch DDP to start distributed training. ElasticFlow uses this launch API to launch distributed training jobs. As specified in §3, developers submit their custom DNN model, hyperparameters of the job, the termination condition, the deadline and other training components like dataset and optimizer to ElasticFlow.

Elastic scaling. ElasticFlow enables elastic scaling in a stop-free manner in multi-node multi-GPU environments. Specifically, when a new scheduling decision is made, ElasticFlow sends the parameters of the running jobs to the workers based on the scheduling decision and then restarts the jobs from the received parameters. The local batch size on each worker is adjusted to maintain the same global batch size. If a running job is suspended, ElasticFlow checkpoints the parameters until it is restarted. To reduce the scaling overhead, ElasticFlow does not delete the CUDA contexts on GPUs and keeps all NCCL process groups active. We evaluate the scaling and migration overheads in §6.6.

Throughput profiling. ElasticFlow makes scheduling decisions based on simulating future job events such as job scaling and job completion. The simulation is done by calculating the number of iterations that can be completed by each job on every job event. To precisely calculate the executed iterations, ElasticFlow first pre-runs each job to profile its throughput with different numbers of GPUs. Then, ElasticFlow profiles its throughput during job execution, and constantly adjust the profiled throughput and the scheduling decisions accordingly. We evaluate the overhead of pre-run profiling in §6.6. The overhead of profiling during job execution is about 0.01s per epoch, which is rather small as most jobs run for hours or even days.

6 EVALUATION

6.1 Methodology

Testbed. Unless specified otherwise, the experiments are conducted on a cluster with 16 servers that includes 128 NVIDIA A100 GPUs in total. Each server is equipped with eight 40GB NVIDIA A100 GPUs, 96 CPU cores, 900 GB RAM, and eight NVIDIA Mellanox HDR InfiniBand HCAs.

Simulator. To evaluate ElasticFlow at larger scales, we develop a simulator using the profiled information in real A100 GPUs. The simulator simulates all job-level events, including job arrival, scaling, and completion. We profile the throughputs of each job with real GPU servers on the testbed as the input of the simulator. To
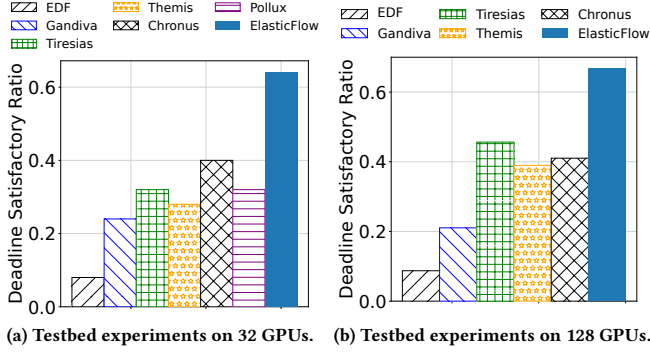
(a) Testbed experiments on 32 GPUs.   (b) Testbed experiments on 128 GPUs.

Figure 6: Deadline satisfactory ratio in testbed experiments.

make the simulator more realistic, we have also measured the job scaling overhead and incorporated it into the simulator. The simulator assigns the overhead to each job on each scheduling event, according to the number of GPUs and the DNN model of the job. Our simulator has very high fidelity, with an error rate of no more than 3% compared with the results in our real cluster experiments.

Workloads. We collect two-month real-world traces from ten clusters in a production system. The cluster size ranges from 164 GPUs to 2,783 GPUs; the number of jobs in each trace ranges from 260 to 15,802. Testbed experiments use traces from one of the clusters; simulations use all the collected traces to evaluate ElasticFlow more comprehensively and at larger scales. Each job in the original traces has the information of submission time, the number of GPUs, and duration. For each job, we randomly choose a DNN model with a batch size from a pool of representative settings listed in Table 1. Similar to previous work [23, 40], we use the duration in the trace and the pre-measured throughput to calculate the number of iterations needed to complete each job. As the traces do not contain any deadline information, we set the deadline for a job at $\lambda \cdot duration$ after its submission, where λ is drawn uniformly from [0.5, 1.5]. λ represents the tightness of the deadline. Note that even when a job's $\lambda < 1$, the cluster may still be able to complete it before the deadline by scaling out the job with more GPUs. For a fair comparison, we also evaluate ElasticFlow with the public Microsoft cluster trace. [25]

Baselines. We compare ElasticFlow to six baselines.

- **Earliest-Deadline-First (EDF):** EDF is a canonical scheduling policy that optimizes for meeting deadlines [17]. It orders jobs based on their deadlines and schedules jobs with the earliest deadline first. It uses as many GPUs as a job can scale out without decreasing the throughput.
- **Gandiva:** Gandiva [58] is a DL scheduler that uses introspective scheduling to refine scheduling decisions continuously. It is not elastic (i.e., use the number of GPUs specified in job traces) and is not deadline-aware.
- **Tiresias:** Tiresias [18] uses two-dimensional scheduling algorithms customized for DL jobs. It is also not elastic and is not deadline-aware.

(a) GPU allocation of different schedulers over time.

(b) Number of submitted and admitted jobs by ElasticFlow over time.

Figure 7: Comparison between different schedulers in testbed experiments.

- **Themis:** Themis [38] provides finish-time fairness for DL jobs. We follow the open-source implementation of Themis in [40]. It is not deadline-aware.
- **Chronus:** Chronus [16] maximizes the number of SLO jobs that can meet deadlines and minimizes the average JCT of best-effort jobs. It is deadline-aware but not elastic.
- **Pollux:** Pollux [48] is a state-of-the-art scheduler for deep learning jobs. It improves the JCT of DL jobs by adaptively co-optimizing statistical efficiency and system throughput. It is elastic but not deadline-aware.

Evaluation metric. The design goal of ElasticFlow is to meet deadlines for DL training jobs. Therefore, the evaluation metric is *deadline satisfactory ratio*, which is the ratio of jobs that can meet their deadlines. We also evaluate *cluster efficiency* (CE), which measures how efficiently the resources in a cluster are being utilized. Let M be the number of GPUs in the cluster and tpt represents throughput, then

$$CE := \frac{1}{M} \sum_{i=1}^{M} \frac{Current\ tpt\ of\ the\ job\ on\ GPU_i}{tpt\ of\ the\ job\ on\ GPU_i\ with\ 1\ GPU}. \quad (8)$$

Because the admission control module of ElasticFlow may drop the jobs with insatiable deadlines, we do not compare the scheduling results on metrics such as job completion time (JCT). We compare the JCT of best-effort jobs when extending ElasticFlow to schedule both SLO and best-effort jobs (§6.5).

6.2 End-to-End Results on Real Testbed

We compare ElasticFlow and the baselines on a real GPU testbed. We follow previous work [58] to speed up the experiments by fast-forwarding. i.e., skips a number of iterations if there are no scheduling events for more than four minutes. The skipped time is calculated by measuring the throughput when the job reaches a

(a) Simulation of all baselines.

(b) Simulation on multiple traces.

Figure 8: Deadline satisfactory ratio under difference traces.

Figure 9: Sources of improvement in ElasticFlow.

**Figure 10: Cluster efficiency of different schedulers.**

(a)   (b)

Figure 11: Performance under a mix of SLO jobs (with deadlines) and best-effort jobs (without deadlines). (a) Deadline satisfactory ratio of SLO jobs. (b) Average JCT of best-effort jobs (normalized to Gandiva).

stable state with the throughput profiling method described in §5. Because Pollux requires the gradients of the DL models' weight during the training process, running DL jobs with Pollux cannot be fast-forwarded. The makespan of running all DL jobs with Pollux depends on the scale of the trace to run. We estimate that using Pollux, running a two-day trace from our collected data on our testbed costs more than $124,000 on Amazon AWS [2]. Due to budget limit, we first compare ElasticFlow and all baselines on a small cluster of 4 servers with 32 GPUs and a small trace with 25 jobs.

Figure 6 (a) shows the results of comparing ElasticFlow with all baselines. ElasticFlow improves the number of jobs that can meet deadlines with EDF, Gandiva, Tiresias, Themis, Chronus, and Pollux to 8.0×, 2.7×, 2.0×, 2.3×, 1.6×, and 2.0×, respectively. ElasticFlow outperforms Gandiva, Tiresias, Themis, and Pollux because ElasticFlow is deadline-aware and drops jobs whose deadlines can not be met. EDF and Chronus are not elastic and cannot fully utilize the GPU resources, leading to low deadline satisfactory ratios.

Then, we compare ElasticFlow with the baselines except for Pollux with a larger scale (on 16 servers with 128 GPUs and a

larger trace with 195 jobs). As is shown in Figure 6(b), ElasticFlow improves the number of jobs that can meet deadlines with EDF, Gandiva, Tiresias, Themis, and Chronus to 7.65 ×, 3.17×, 1.46×, 1.71×, and 1.62×, respectively. We will discuss the sources of improvement of ElasticFlow in § 6.1

Figure 7 shows the number of allocated GPUs and the number of submitted and admitted jobs during the execution on the testbed. From the figure, we observe that when there are few jobs submitted to the cluster, the resource contention in the cluster is low. ElasticFlow can take full advantage of the idle GPU resources so that admitted jobs can complete earlier. Schedulers such as Gandiva and Tiresias are not elastic and thus cannot utilize idle GPUs. When there is a burst of job submissions (e.g., the 13[th] hour), some jobs are dropped to guarantee the deadlines of admitted jobs.

6.3 End-to-End Results in Simulations

We compare ElasticFlow and the baselines with simulation on more settings. The simulation of Pollux requires to profile step time as well as other detailed training statuses of thousands of DL model settings and topology settings, and the estimated cost is about $90,000 on Amazon AWS. We do not profile them on our testbed due to budget limit. Instead, we use the profiled data provided by Pollux and transform the 195-job trace used in § 6.2 into the form with which we can run the Pollux simulation. The data provided
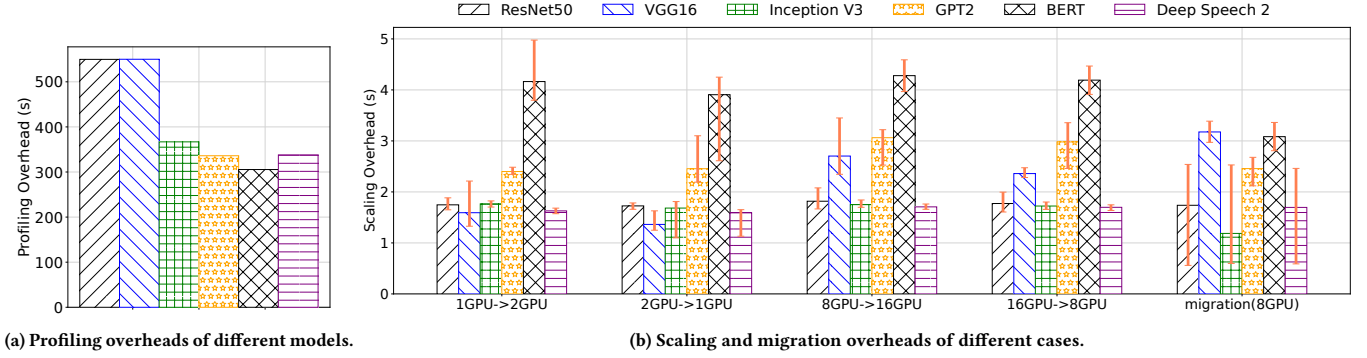
(a) Profiling overheads of different models.   (b) Scaling and migration overheads of different cases.

Figure 12: System overheads of ElasticFlow.

by Pollux was profiled on a cluster of 16 servers with 64 NVIDIA T4 GPUs. Figure 8(a) shows the simulation results. Similar to the results of testbed experiments, ElasticFlow achieves a much higher deadline satisfactory ratio than all baselines.

Figure 8(b) compares ElasticFlow with the baselines except for Pollux on more traces and larger scales. We use the job traces collected from ten different clusters and one public trace of the Microsoft Philly cluster [25]. We observe that ElasticFlow consistently outperforms the baselines in all traces. The deadline satisfactory ratios of Gandiva and Tiresias do not change much across the traces because they are not deadline-aware and lack elasticity. In some traces (trace #9 and #10), EDF achieves a higher deadline satisfactory ratio than Gandiva, Tiresias, Themis, and Chronus. This is because when the cluster size is large enough for the jobs in these traces, EDF runs jobs with the earliest deadline first and uses adequate GPU resources to finish them as quickly as possible. However, EDF does not make the most efficient use of GPUs due to the sub-linear scaling of DL training jobs. Therefore, in other traces where the GPU resources in the cluster are not adequate, EDF performs poorly. On average, ElasticFlow improves the deadline satisfactory ratio to 12.95×, 2.58×, 2.15×, 1.76×, and 1.68× compared to EDF, Gandiva, Tiresias, Themis, and Chronus, respectively.

6.4 Sources of Improvement in ElasticFlow

Improvement in deadline satisfactory ratio. We analyze the sources of improvement in ElasticFlow. ElasticFlow has two key components, which are admission control and elastic scaling. To show that both components are important, we develop two variants on top of EDF, which are EDF + Admission Control and EDF + Elastic Scaling. We vary the cluster size and keep the same load. First, Figure 9 shows that both admission control and elastic scaling contribute to the improvement of ElasticFlow, as admission control drops the jobs with unsatisfiable deadlines and elastic scaling adjusts the resource allocation in time according to cluster status and the jobs' deadlines. Only adding one of them to EDF performs worse than ElasticFlow. Second, we observe that the gap between EDF + Elastic Scaling and ElasticFlow decreases when the cluster size increases. This is because when the load is kept the same, most jobs can be admitted with more GPUs, and elastic scaling is critical to the improvement. On the other hand, when the cluster size is

small, admission control is important to avoid wasting resources on jobs that can never meet their deadlines.

Improvement in cluster efficiency. We evaluate how well ElasticFlow utilizes the cluster resources. We simulate the execution of a 100-job trace on a 16-node cluster where each node has 8 GPUs. If ElasticFlow declines any job in the trace, the comparison is unfair to the baselines. Therefore, we set the deadlines to be loose enough to admit all jobs (1.5× of a job's duration). This ensures that all solutions run the same set of jobs. Figure 10 shows that ElasticFlow achieves a higher cluster efficiency than the baselines in the first 100 hours to save more GPU resources for later jobs. This is because the resource allocation module allocates the idle GPUs in the cluster in a way that the GPUs are used efficiently. ElasticFlow also achieves the smallest makespan when finishing all jobs.

6.5 Handling Jobs without Deadlines

ElasticFlow can schedule both SLO jobs and best-efforts jobs. We vary the percentage of best-effort jobs and measure the deadline satisfactory ratio of SLO jobs and the average JCT of best-effort jobs. Because the average JCT under EDF is too large compared to other solutions, we normalize the average JCT of each solution to that of Gandiva for illustration. As is shown in Figure 11, ElasticFlow achieves the highest deadline satisfactory ratio for jobs with deadline requirements among all the solutions. For traces with 10% of best-effort jobs, ElasticFlow achieves the smallest average JCT for best-effort jobs. In other traces, ElasticFlow reserves more GPU resources for jobs with deadline requirements and sacrifices JCT for deadline satisfactory ratio as expected.

6.6 System Overheads

Profiling overheads. Figure 12(a) shows the profiling overheads of different DL models. ElasticFlow profiles the throughput of each new DL model with different numbers of GPUs along with different batch sizes. ElasticFlow records the largest local batch size of each job that the GPU memory can hold, and records the largest and smallest number of GPUs for each job to avoid poor performance or memory overflow. If adding more GPU to a job with a specific batch size cannot increase the throughput of the given job, ElasticFlow stops the profiling procedure for that batch size. As DL training

usually takes hours, days or even weeks, the profiling overhead is marginal. There is no need to profile known/repeated jobs.

Scaling and migration overheads. The scaling and migration overhead is the interval of "*suspending a job*" and "*restarting the job on a new set of GPUs*". Figure 12(b) shows the scaling and migration overheads of different DL models for five cases. The first four cases are changing the number of GPUs for a job, and the fifth case is to change the set of GPUs (i.e., migrating to 8 GPUs on another machine). Our current prototype uses simple checkpointing/restore available in PyTorch for scaling and migration. When the number of GPUs is large, the overheads of the internal implementation of checkpointing/restore in PyTorch dominate. When we scale up or down from a single GPU, the ElasticFlow overheads are visible. Therefore, the overheads of different cases are similar; the variance is most from the internal implementation of checkpointing/restore in PyTorch. This implementation can be further optimized by keeping the model weights in memory if the GPU devices will train the same DL model after a scaling event. The average scheduling interval of ElasticFlow on the real-world traces we collected in a production system is about 23 minutes. Compared to the scheduling interval, the scaling and migration overheads are marginal.

7 CONCLUSION

In this paper, we proposed ElasticFlow, an elastic serverless computing platform for distributed training. ElasticFlow performs admission control based on minimum satisfactory share to guarantee deadlines. We developed a scheduling algorithm that dynamically allocates resources to admitted jobs based on diminishing returns. ElasticFlow applies buddy allocation to job placement to eliminate the effect of topology. The evaluation showed that ElasticFlow increased the number of jobs that can meet their deadlines by 1.46–7.65× over existing solutions. This work does not raise any ethical issues.

REFERENCES

[1] 2019. NCCL. https://developer.nvidia.com/nccl. Retrieved on July 3, 2022.
[2] 2021. AWS. https://aws.amazon.com/console/?nc1=h_ls. Retrieved on July 3, 2022.
[3] 2022. Amazon SageMaker. https://aws.amazon.com/cn/sagemaker/. Retrieved on July 3, 2022.
[4] 2022. gRPC. https://grpc.io. Retrieved on July 3, 2022.

[5] 2022. TorchElastic. https://pytorch.org/elastic/0.2.0/index.html. Retrieved on July 3, 2022.
[6] Dario Amodei, Sundaram Ananthanarayanan, Rishita Anubhai, Jingliang Bai, Eric Battenberg, Carl Case, Jared Casper, Bryan Catanzaro, Jingdong Chen, Mike Chrzanowski, Adam Coates, Greg Diamos, Erich Elsen, Jesse H. Engel, Linxi Fan, Christopher Fougner, Awni Y. Hannun, Billy Jun, Tony Han, Patrick LeGresley, Xiangang Li, Libby Lin, Sharan Narang, Andrew Y. Ng, Sherjil Ozair, Ryan Prenger, Sheng Qian, Jonathan Raiman, Sanjeev Satheesh, David Seetapun, Shubho Sengupta, Chong Wang, Yi Wang, Zhiqian Wang, Bo Xiao, Yan Xie, Dani Yogatama, Jun Zhan, and Zhenyao Zhu. 2016. Deep Speech 2 : End-to-End Speech Recognition in English and Mandarin. In *Proceedings of the 33nd International Conference on Machine Learning, ICML 2016*, Vol. 48. 173–182.
[7] Sanjith Athlur, Nitika Saran, Muthian Sivathanu, Ramachandran Ramjee, and Nipun Kwatra. 2022. Varuna: scalable, low-cost training of massive deep learning models. In *Proceedings of 17th European Conference on Computer Systems, EuroSys 2022*. 472–487.
[8] Dimitri P. Bertsekas and Robert G. Gallager. 1992. *Data Networks, Second Edition*. Prentice Hall.
[9] Scott Boag, Parijat Dube, Benjamin Herta, Waldemar Hummer, Vatche Ishakian, K Jayaram, Michael Kalantar, Vinod Muthusamy, Priya Nagpurkar, and Florian Rosenberg. 2017. Scalable Multi-framework Multi-tenant Lifecycle Management of Deep Learning Training Jobs. In *Workshop on ML Systems, NeurIPS 2017*.
[10] Shubham Chaudhary, Ramachandran Ramjee, Muthian Sivathanu, Nipun Kwatra, and Srinidhi Viswanatha. 2020. Balancing efficiency and fairness in heterogeneous GPU clusters for deep learning. In *Proceedings of the 15th European Conference on Computer Systems, EuroSys 2020*. 1:1–1:16.
[11] Yangrui Chen, Yanghua Peng, Yixin Bao, Chuan Wu, Yibo Zhu, and Chuanxiong Guo. 2020. Elastic parameter server load distribution in deep learning clusters. In *Proceedings of the 11th ACM Symposium on Cloud Computing, SoCC 2020*. 507–521.
[12] Jeffrey Dean, Greg Corrado, Rajat Monga, Kai Chen, Matthieu Devin, Quoc V. Le, Mark Z. Mao, Marc'Aurelio Ranzato, Andrew W. Senior, Paul A. Tucker, Ke Yang, and Andrew Y. Ng. 2012. Large Scale Distributed Deep Networks. In *Proceedings of 26th Annual Conference on Neural Information Processing Systems, NeurIPS 2012*. 1232–1240.
[13] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Fei-Fei Li. 2009. ImageNet: A Large-scale Hierarchical Image Database. In *Proceedings of 2009 IEEE Computer Society Conference on Computer Vision and Pattern Recognition, CVPR 2009*. 248–255.
[14] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Proceedings of 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019*. 4171–4186.
[15] Anis Elgabli, Jihong Park, Amrit S Bedi, Mehdi Bennis, and Vaneet Aggarwal. 2020. GADMM: Fast and Communication Efficient Framework for Distributed Machine Learning. *Journal of Machine Learning Research* 21, 76 (2020), 1–39.
[16] Wei Gao, Zhisheng Ye, Peng Sun, Yonggang Wen, and Tianwei Zhang. 2021. Chronus: A Novel Deadline-aware Scheduler for Deep Learning Training Jobs. In *Proceedings of the 12th ACM Symposium on Cloud Computing, Seattle, SoCC 2021*. 609–623.
[17] Laurent George, Nicolas Rivierre, and Marco Spuri. 1996. *Preemptive and non-preemptive real-time uniprocessor scheduling*. Ph. D. Dissertation. Inria.
[18] Juncheng Gu, Mosharaf Chowdhury, Kang G. Shin, Yibo Zhu, Myeongjae Jeon, Junjie Qian, Hongqiang Harry Liu, and Chuanxiong Guo. 2019. Tiresias: A GPU Cluster Manager for Distributed Deep Learning. In *Proceedings of 16th USENIX Symposium on Networked Systems Design and Implementation, NSDI 2019*. 485–500.
[19] Aaron Harlap, Alexey Tumanov, Andrew Chung, Gregory R Ganger, and Phillip B Gibbons. 2017. Proteus: Agile ML Elasticity through Tiered Reliability in Dynamic Resource Markets. In *Proceedings of the 12th European Conference on Computer Systems, EuroSys 2017*. 589–604.
[20] Sayed Hadi Hashemi, Sangeetha Abdu Jyothi, and Roy H. Campbell. 2019. TicTac: Accelerating Distributed Deep Learning with Communication Scheduling. In *Proceedings of Machine Learning and Systems 2019, MLSys 2019*.
[21] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition, CVPR 2016*. 770–778.
[22] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Dehao Chen, Mia Xu Chen, HyoukJoong Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng Chen. 2019. GPipe: Efficient Training of Giant Neural Networks using Pipeline Parallelism. In *Proceedings of 33th Annual Conference on Neural Information Processing Systems, NeurIPS 2019*. 103–112.
[23] Changho Hwang, Taehyun Kim, Sunghyun Kim, Jinwoo Shin, and KyoungSoo Park. 2021. Elastic Resource Sharing for Distributed Deep Learning. In *Proceedings of 18th USENIX Symposium on Networked Systems Design and Implementation, NSDI 2021*. 721–739.
[24] Abhinav Jangda, Jun Huang, Guodong Liu, Amir Hossein Nodehi Sabet, Saeed Maleki, Youshan Miao, Madanlal Musuvathi, Todd Mytkowicz, and Olli Saarikivi.

2022. Breaking the computation and communication abstraction barrier in distributed machine learning workloads. In *Proceedings of 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS 2022.* 402–416.

[25] Myeongjae Jeon, Shivaram Venkataraman, Amar Phanishayee, Junjie Qian, Wencong Xiao, and Fan Yang. 2019. Analysis of Large-Scale Multi-Tenant GPU Clusters for DNN Training Workloads. In *Proceedings of 2019 USENIX Annual Technical Conference, ATC 2019.* 947–960.

[26] Myeongjae Jeon, Shivaram Venkataraman, Junjie Qian, Amar Phanishayee, Wencong Xiao, and Fan Yang. 2018. Multi-tenant GPU clusters for deep learning workloads: Analysis and implications. *Technical report, Microsoft Research* (2018).

[27] Zhihao Jia, Oded Padon, James Thomas, Todd Warszawski, Matei Zaharia, and Alex Aiken. 2019. TASO: Optimizing Deep Learning Computation with Automatic Generation of Graph Substitutions. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP 2019.* 47–62.

[28] Zhihao Jia, Matei Zaharia, and Alex Aiken. 2019. Beyond Data and Model Parallelism for Deep Neural Networks. In *Proceedings of Machine Learning and Systems 2019, MLSys 2019.*

[29] Yimin Jiang, Yibo Zhu, Chang Lan, Bairen Yi, Yong Cui, and Chuanxiong Guo. 2020. A Unified Architecture for Accelerating Distributed DNN Training in Heterogeneous GPU/CPU Clusters. In *Proceedings of 14th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2020.* 463–479.

[30] Eric Jonas, Johann Schleier-Smith, Vikram Sreekanti, Chia-che Tsai, Anurag Khandelwal, Qifan Pu, Vaishaal Shankar, Joao Carreira, Karl Krauth, Neeraja Jayant Yadwadkar, Joseph E. Gonzalez, Raluca Ada Popa, Ion Stoica, and David A. Patterson. 2019. Cloud Programming Simplified: A Berkeley View on Serverless Computing. *arXiv preprint arXiv:1902.03383* (2019).

[31] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. 2012. ImageNet Classification with Deep Convolutional Neural Networks. In *Proceedings of 26th Annual Conference on Neural Information Processing Systems, NeurIPS 2012.* 1106–1114.

[32] Mu Li, David G Andersen, Jun Woo Park, Alexander J Smola, Amr Ahmed, Vanja Josifovski, James Long, Eugene J Shekita, and Bor-Yiing Su. 2014. Scaling Distributed Machine Learning with the Parameter Server. In *Proceedings of 11th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2014.* 583–598.

[33] Mu Li, David G. Andersen, Alexander J. Smola, and Kai Yu. 2014. Communication Efficient Distributed Machine Learning with the Parameter Server. In *Proceedings of 28th Annual Conference on Neural Information Processing Systems, NeurIPS 2014.* 19–27.

[34] Richard Liaw, Romil Bhardwaj, Lisa Dunlap, Yitian Zou, Joseph E. Gonzalez, Ion Stoica, and Alexey Tumanov. 2019. HyperSched: Dynamic Resource Reallocation for Model Development on a Deadline. In *Proceedings of the ACM Symposium on Cloud Computing, SoCC 2019.* 61–73.

[35] Qinyi Luo, Jiaao He, Youwei Zhuo, and Xuehai Qian. 2020. Prague: High-Performance Heterogeneity-Aware Asynchronous Decentralized Training. In *Architectural Support for Programming Languages and Operating Systems, ASPLOS 2020.* 401–416.

[36] Qinyi Luo, Jinkun Lin, Youwei Zhuo, and Xuehai Qian. 2019. Hop: Heterogeneity-aware Decentralized Training. In *Proceedings of the 24th International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS 2019.* 893–907.

[37] Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher Potts. 2011. Learning Word Vectors for Sentiment Analysis. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies, ACL 2022.* 142–150.

[38] Kshiteej Mahajan, Arjun Balasubramanian, Arjun Singhvi, Shivaram Venkataraman, Aditya Akella, Amar Phanishayee, and Shuchi Chawla. 2020. Themis: Fair and Efficient GPU Cluster Scheduling. In *Proceedings of 17th USENIX Symposium on Networked Systems Design and Implementation, NSDI 2020.* 289–304.

[39] Luo Mai, Guo Li, Marcel Wagenländer, Konstantinos Fertakis, Andrei-Octavian Brabete, and Peter R. Pietzuch. 2020. KungFu: Making Training in Distributed Machine Learning Adaptive. In *Proceedings of 14th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2020.* 937–954.

[40] Deepak Narayanan, Keshav Santhanam, Fiodar Kazhamiaka, Amar Phanishayee, and Matei Zaharia. 2020. Heterogeneity-Aware Cluster Scheduling Policies for Deep Learning Workloads. In *Proceedings of 14th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2020.* 481–498.

[41] Andrew Or, Haoyu Zhang, and Michael J. Freedman. 2020. Resource Elasticity in Distributed Deep Learning. In *Proceedings of Machine Learning and Systems 2020, MLSys 2020.*

[42] Andrew Or, Haoyu Zhang, and Michael J. Freedman. 2020. VirtualFlow: Decoupling Deep Learning Model Execution from Underlying Hardware. *arXiv preprint arXiv:2009.09523* (2020).

[43] Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. 2015. Librispeech: An ASR corpus based on public domain audio books. In *Proceedings of 2015 IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2015.* 5206–5210.

[44] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. 2019. Pytorch: An imperative style, high-performance deep learning library. *arXiv preprint arXiv:1912.01703* (2019).

[45] Yanghua Peng, Yixin Bao, Yangrui Chen, Chuan Wu, and Chuanxiong Guo. 2018. Optimus: an Efficient Dynamic Resource Scheduler for Deep Learning Clusters. In *Proceedings of the 13th European Conference on Computer Systems, EuroSys 2018.* 1–14.

[46] Yanghua Peng, Yixin Bao, Yangrui Chen, Chuan Wu, Chen Meng, and Wei Lin. 2019. DL2: A Deep Learning-driven Scheduler for Deep Learning Clusters. *arXiv preprint arXiv:1909.06040* (2019).

[47] Aurick Qiao, Abutalib Aghayev, Weiren Yu, Haoyang Chen, Qirong Ho, Garth A. Gibson, and Eric P. Xing. 2018. Litz: Elastic Framework for High-Performance Distributed Machine Learning. In *Proceedings of 2018 USENIX Annual Technical Conference, ATC 2018.* 631–644.

[48] Aurick Qiao, Sang Keun Choe, Suhas Jayaram Subramanya, Willie Neiswanger, Qirong Ho, Hao Zhang, Gregory R. Ganger, and Eric P. Xing. 2021. Pollux: Co-adaptive Cluster Scheduling for Goodput-Optimized Deep Learning. In *Proceedings of 15th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2021.* 1–18.

[49] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners. *OpenAI blog* 1, 8 (2019), 9.

[50] John E. Shore. 1975. On the External Storage Fragmentation Produced by First-Fit and Best-Fit Allocation Strategies. *Commun. ACM* 18, 8 (1975), 433–440.

[51] Karen Simonyan and Andrew Zisserman. 2015. Very Deep Convolutional Networks for Large-Scale Image Recognition. In *Proceedings of 3rd International Conference on Learning Representations, ICLR 2015.*

[52] Muthian Sivathanu, Tapan Chugh, Sanjay S Singapuram, and Lidong Zhou. 2019. Astra: Exploiting Predictability to Optimize Deep Learning. In *Proceedings of the 24th International Conference on Architectural Support for Programming Languages and Operating Systems, ASPLOS 2019.* 909–923.

[53] Suhas Jayaram Subramanya, Harsha Vardhan Simhadri, Srajan Garg, Anil Kag, and Venkatesh Balasubramanian. 2019. BLAS-on-flash: An Efficient Alternative for Large Scale ML Training and Inference?. In *Proceedings of 16th USENIX Symposium on Networked Systems Design and Implementation, NSDI 2019.* 469–484.

[54] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jonathon Shlens, and Zbigniew Wojna. 2016. Rethinking the Inception Architecture for Computer Vision. In *Proceedings of 2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016.* 2818–2826.

[55] Jianyu Wang and Gauri Joshi. 2018. Adaptive communication strategies to achieve the best error-runtime trade-off in local-update SGD. *arXiv preprint arXiv:1810.08313* (2018).

[56] Alex Warstadt, Amanpreet Singh, and Samuel R Bowman. 2018. Neural Network Acceptability Judgments. *arXiv preprint arXiv:1805.12471* (2018).

[57] Yidi Wu, Kaihao Ma, Xiao Yan, Zhi Liu, Zhenkun Cai, Yuzhen Huang, James Cheng, Han Yuan, and Fan Yu. 2019. Elastic Deep Learning in Multi-tenant GPU Cluster. *arXiv preprint arXiv:1909.11985* (2019).

[58] Wencong Xiao, Romil Bhardwaj, Ramachandran Ramjee, Muthian Sivathanu, Nipun Kwatra, Zhenhua Han, Pratyush Patel, Xuan Peng, Hanyu Zhao, Quanlu Zhang, Fan Yang, and Lidong Zhou. 2018. Gandiva: Introspective Cluster Scheduling for Deep Learning. In *Proceedings of 13th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2018.* 595–610.

[59] Lei Xie, Jidong Zhai, Baodong Wu, Yuanbo Wang, Xingcheng Zhang, Peng Sun, and Shengen Yan. 2020. Elan: Towards Generic and Efficient Elastic Training for Deep Learning. In *Proceedings of 40th IEEE International Conference on Distributed Computing Systems, ICDCS 2020.*

[60] Hao Zhang, Zeyu Zheng, Shizhen Xu, Wei Dai, Qirong Ho, Xiaodan Liang, Zhiting Hu, Jinliang Wei, Pengtao Xie, and Eric P Xing. 2017. Poseidon: An Efficient Communication Architecture for Distributed Deep Learning on GPU Clusters. In *Proceedings of 2017 USENIX Annual Technical Conference, ATC 2017.* 181–193.

[61] Xin Zhang, Jia Liu, Zhengyuan Zhu, and Elizabeth S Bentley. 2019. Compressed Distributed Gradient Descent: Communication-efficient Consensus over Networks. In *IEEE INFOCOM 2019-IEEE Conference on Computer Communications.* IEEE, 2431–2439.

[62] Hanyu Zhao, Zhenhua Han, Zhi Yang, Quanlu Zhang, Fan Yang, Lidong Zhou, Mao Yang, Francis C. M. Lau, Yuqi Wang, Yifan Xiong, and Bin Wang. 2020. HiveD: Sharing a GPU Cluster for Deep Learning with Guarantees. In *Proceedings of 14th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2020.* 515–532.

[63] Yiren Zhao, Ilia Shumailov, Robert D. Mullins, and Ross Anderson. 2019. To Compress Or Not To Compress: Understanding The Interactions Between Adversarial Attacks And Neural Network Compression. In *Proceedings of Machine Learning and Systems 2019, MLSys 2019.*